

Topic Modeling with Mallet

Sowmya

3/20/2018

Purpose of this tutorial is to expand on what the author describes in Chapter 13. You have to first read Chapters 10 and 13 (and make notes), before trying this out.

Software Requirements:

Java needs to be installed on your computer. Go to this link for installation: https://java.com/en/download/help/index_installing.xml

After that: R library `mallet` needs to be installed in Rstudio (which inturn installs `rJava`). On MacOS, R may also prompt you to install some other legacy java version - if it does, install that.

R libraries you need to install: `mallet`, `wordcloud`, `XML`

Data pre-processing:

The author again uses the XML corpus from Text classification example - this is provided as supplementary material in the electronic version of the textbook. I am providing it in the data folder inside this folder.

He first splits each file into 1000 chunks (instead of 10 in the classification example) using the function below and does these pre-processing: Extract text from xml, lowercase it, replace anything that is not an alpha-numeric character or a space or an apostrophe with a space.

Then, the next part is about splitting the text into multiple parts. Two methods are suggested for that. First one is - split the text into X parts based on the length of the text i.e., split after every 1000 words (for example). Second method is: splitting all texts into equal number of words. If you do that, sometimes, there may be a few leftover words at the end. This function addresses that too (how?).

```
makeFlexTextChunks <- function(doc.object, chunk.size=1000, percentage=TRUE){
  paras <- getNodeSet(doc.object,
    "/d:TEI/d:text/d:body//d:p",
    c(d = "http://www.tei-c.org/ns/1.0"))
  words <- paste(sapply(paras,xmlValue), collapse=" ")
  words.lower <- tolower(words)
  words.lower <- gsub("[^[:alnum:][:space:]]'", " ", words.lower)
  words.l <- strsplit(words.lower, "\\s+")
  word.v <- unlist(words.l)
  x <- seq_along(word.v)
  if(percentage){
    max.length <- length(word.v)/chunk.size
    chunks.l <- split(word.v, ceiling(x/max.length))
  } else {
    chunks.l <- split(word.v, ceiling(x/chunk.size))
    #deal with small chunks at the end
    if(length(chunks.l[[length(chunks.l)]]) <=
      length(chunks.l[[length(chunks.l)]]/2){
      chunks.l[[length(chunks.l)-1]] <-
        c(chunks.l[[length(chunks.l)-1]],
          chunks.l[[length(chunks.l)]])
    }
  }
}
```

```

    chunks.l[[length(chunks.l)]] <- NULL
  }
}
chunks.l <- lapply(chunks.l, paste, collapse=" ")
chunks.df <- do.call(rbind, chunks.l)
return(chunks.df)
}

```

Once you have that function, the next step is to use this function to create such chunks for each file, combine them into a two column data frame where one column represents the text and chunk id and the other column shows the words in that chunk.

```

library(XML)
wd <- "~/Dropbox/ClassroomSlides-BothCourses/LING410X/Week11Mats/21Mar2018/"
inputDir <- paste(wd, "data/XMLAuthorCorpus", sep="")
stopwords_path <- paste(wd, "data/stoplist.csv", sep="")
files.v <- dir(path=inputDir, pattern="*.xml")
chunk.size <- 1000
topic.m <- NULL
for(i in 1:length(files.v)){
  doc.object <- xmlTreeParse(file.path(inputDir, files.v[i]),
                             useInternalNodes=TRUE)
  chunk.m <- makeFlexTextChunks(doc.object, chunk.size,
                                percentage=FALSE)
  textname <- gsub("\\.*", "", files.v[i])
  segments.m <- cbind(paste(textname,
                             segment=1:nrow(chunk.m), sep="_"), chunk.m)
  topic.m <- rbind(topic.m, segments.m)
}
documents <- as.data.frame(topic.m, stringsAsFactors=F)
colnames(documents) <- c("id", "text")

```

This is all the pre-processing needed before you start training a topic model. `mallet.import` takes 5 arguments - ids for documents, text of the documents, stopwords list (csv file), preserve case (true or false), and a regular expression to split the text into tokens/words.

Once you have these, creating a topic model is a simple procedure. You need to specify the number of topics.

```

library(mallet)

## Loading required package: rJava

mallet.instances <- mallet.import(documents$id, documents$text, stopwords_path, FALSE,
                                  token.regex="([\\p{L}']+)")
topic.model <- MalletLDA(num.topics=43)
topic.model$loadDocuments(mallet.instances)

#The following line is optional.
#Mallet takes default parameters if you don't specify anything here.
topic.model$setAlphaOptimization(40, 80)

#This starts the training process
topic.model$train(400)

```

We can use `explore` this topic model to analyse topics and words in this data:

[illegible]

```
vocabulary <- topic.model$getVocabulary()
```

```
length(vocabulary)
```

```
## [1] 55444
```

```
head(vocabulary)
```

```
## [1] "summer"    "topsail"   "schooner"  "slipped"   "cove"      "trinidad"
```

```
vocabulary[1:50]
```

##	[1]	"summer"	"topsail"	"schooner"	"slipped"	"cove"
##	[6]	"trinidad"	"head"	"dropped"	"anchor"	"edge"
##	[11]	"kelp"	"fields"	"fifteen"	"minutes"	"small"
##	[16]	"boat"	"deposited"	"beach"	"man"	"armed"
##	[21]	"long"	"squirrel"	"rifle"	"axe"	"carrying"
##	[26]	"food"	"clothing"	"brown"	"canvas"	"pack"
##	[31]	"watched"	"return"	"weigh"	"stand"	"sea"
##	[36]	"northwest"	"trades"	"disappeared"	"ken"	"swung"
##	[41]	"broad"	"powerful"	"back"	"strode"	"resolutely"
##	[46]	"timber"	"mouth"	"river"	"john"	"cardigan"

#word freqs has 3 columns. word, its frequency in entire corpus, its frequency per document.

```
word.freqs <- mallet.word.freqs(topic.model)
```

#this below line returns a matrix where each row is a topic,

#each column is a word in the dataset.

```
topic.words.m <- mallet.topic.words(topic.model,smoothed=TRUE,normalized=TRUE)
```

```
dim(topic.words.m)
```

```
## [1] 43 55444
```

```
rowSums(topic.words.m)
```

```
## [1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

```
## [36] 1 1 1 1 1 1 1 1
```

```
topic.words.m[1:3, 1:3]
```

```
##           [,1]      [,2]      [,3]
```

```
## [1,] 3.440089e-07 3.440089e-07 3.440089e-07
```

```
## [2.] 8.585185e-07 8.585185e-07 8.585185e-07
```

```
## [3.] 7.240594e-07 7.240594e-07 7.240594e-07
```

```
colnames(topic.words.m) <- vocabulary
```

```
keywords <- c("california", "ireland")
```

```
topic.words.m[, keywords]
```

```
##      california      ireland
```

```
## [1,] 3.440089e-07 3.440089e-07
```

```
## [2,] 8.585185e-07 8.585185e-07
```

```
## [3.] 7.240594e-07 7.240594e-07
```

```
## [4.] 5.918609e-07 5.918609e-07
```

```
## [5.] 2.315544e-07 2.315544e-07
```

```
## [6,] 1.337125e-07 1.337125e-07
```

```
## [7.] 3.354310e-07 3.354310e-07
```

```
## [8,] 1.937672e-07 1.937672e-07
## [9,] 1.517795e-07 1.517795e-07
## [10,] 1.632222e-07 1.632222e-07
## [11,] 3.363235e-03 5.582788e-07
## [12,] 5.113770e-04 7.441480e-07
## [13,] 8.708564e-04 8.150889e-07
## [14,] 1.669692e-06 1.669692e-06
## [15,] 5.532627e-07 5.532627e-07
## [16,] 1.340916e-06 1.340916e-06
## [17,] 6.742904e-07 6.742904e-07
## [18,] 9.597928e-07 1.171818e-03
## [19,] 1.738793e-07 1.738793e-07
## [20,] 4.354327e-07 4.354327e-07
## [21,] 1.918428e-07 1.918428e-07
## [22,] 8.544914e-07 8.544914e-07
## [23,] 2.621855e-03 3.800916e-03
## [24,] 3.747765e-07 3.747765e-07
## [25,] 1.326942e-06 1.326942e-06
## [26,] 1.868290e-06 1.141439e-03
## [27,] 3.153644e-07 3.153644e-07
## [28,] 1.221412e-06 1.221412e-06
## [29,] 2.484411e-07 2.484411e-07
## [30,] 1.470912e-04 1.100059e-03
## [31,] 3.644561e-07 3.644561e-07
## [32,] 2.880987e-07 4.591153e-03
## [33,] 1.570525e-06 1.570525e-06
## [34,] 3.203872e-07 3.203872e-07
## [35,] 8.723949e-07 8.723949e-07
## [36,] 9.387479e-04 4.922371e-07
## [37,] 8.205965e-07 8.205965e-07
## [38,] 2.280596e-07 2.280596e-07
## [39,] 6.784761e-07 6.784761e-07
## [40,] 3.130012e-07 5.491972e-04
## [41,] 7.903345e-07 1.458333e-02
## [42,] 7.749249e-07 7.749249e-07
## [43,] 1.871218e-02 7.597964e-07
```

```
#The topic which has the highest count of numbers in the row (indicating weights for words)
#can be thought of as an important topic.
imp.row <- which(rowSums(topic.words.m[, keywords]) == max(rowSums(topic.words.m[, keywords])))
```

```
#10 most frequent words in the Topic given by imp.row
mallet.top.words(topic.model, topic.words.m[imp.row,], 10)
```

```
##          words      weights
## san          san 0.020392147
## california  california 0.018712175
## men          men 0.015989461
## mr           mr 0.013846049
## land         land 0.013150888
## francisco    francisco 0.013035027
## city         city 0.012687447
## state        state 0.010022663
## states       states 0.007821320
## gold         gold 0.006199278
```

```
mallet.top.words(topic.model, topic.words.m[10,], 5)
```

##	words	weights
##	god	god 0.01987442
##	heart	heart 0.01932686
##	father	father 0.01884151
##	mother	mother 0.01492141
##	poor	poor 0.01273114

Little bit of visualization

This needs the library wordcloud.

```
library(wordcloud)
```

```
## Loading required package: RColorBrewer
```

```
#plotting a word cloud for the imp.row
```

```
topic.top.words <- mallet.top.words(topic.model,topic.words.m[imp.row,], 100)
```

```
wordcloud(topic.top.words$words,topic.top.words$weights,c(4,.8), rot.per=0, random.order=F)
```



#plotting a word cloud for topic 20

```
topic.top.words <- mallet.top.words(topic.model,topic.words.m[20,], 100)
```

```
wordcloud(topic.top.words$words,topic.top.words$weights,c(4,.8), rot.per=0, random.order=F)
```

persons delightful
agreeable admiration amused
stage english interesting
stories charming bow flowers
tall pleased company ball carriage
stranger large dinner made party french
dancing play wine ladies handsome mr girls
sex friends half beauty crowd
man hall dance miss set red lost
supper side gave music house happy
making high young fine laughing
praised called lady time lord general
sweet jack great words grand
grace laugh table pretty beautiful air pleasure
attention women fond conversation good gay merry year
light singing manners speech round art drawing
entered laughter pleasant small
slight dozen gentleman generally
acquaintance daughter learned