

# LING 410X: Language as Data

Semester: Spring '18

Instructor: Sowmya Vajjala

Iowa State University, USA

13 February 2018

# Class Outline

- ▶ Recap of last week and how to continue on those topics
- ▶ Searching for words and their context in text.
- ▶ Getting ngrams from text
- ▶ Assignment 3 description

Note: Assignment 2 discussion on Thursday, not Today.

# Recap

# Recap of last week

- ▶ Text analysis concepts:
  - ▶ Lexical variety: mean word frequency, type-token ratio, hapax richness (percentage of words that appear only once)

# Recap of last week

- ▶ Text analysis concepts:
  - ▶ Lexical variety: mean word frequency, type-token ratio, hapax richness (percentage of words that appear only once)
- ▶ R concepts:
  1. Data structures in R: vectors, lists, matrices, arrays, dataframes, factors

# Recap of last week

- ▶ Text analysis concepts:
  - ▶ Lexical variety: mean word frequency, type-token ratio, hapax richness (percentage of words that appear only once)
- ▶ R concepts:
  1. Data structures in R: vectors, lists, matrices, arrays, dataframes, factors
  2. Usage of built-in R functions such as: lapply, unname, sum, mean, scale

# Recap of last week

- ▶ Text analysis concepts:
  - ▶ Lexical variety: mean word frequency, type-token ratio, hapax richness (percentage of words that appear only once)
- ▶ R concepts:
  1. Data structures in R: vectors, lists, matrices, arrays, dataframes, factors
  2. Usage of built-in R functions such as: lapply, unname, sum, mean, scale
  3. Writing our own R functions

# Recap of last week

- ▶ Text analysis concepts:
  - ▶ Lexical variety: mean word frequency, type-token ratio, hapax richness (percentage of words that appear only once)
- ▶ R concepts:
  1. Data structures in R: vectors, lists, matrices, arrays, dataframes, factors
  2. Usage of built-in R functions such as: lapply, unname, sum, mean, scale
  3. Writing our own R functions
- ▶ Reporting analysis process and results using R Markdown



# How to work further on these

- ▶ I covered Chapters 6 and 7 from the text book during last week. Go through those chapters, do the exercises at the end of the chapter.
- ▶ Revisit the slides, and understand the lines of R code shown in class.
- ▶ Try to do it yourself, taking some other text file or
- ▶ converting the programs to work on a collection of files instead of one file (How??)
- ▶ Do the lab exercises whenever you find time, if you did not finish them in lab.
- ▶ Think in terms of functions you can create and reuse for the rest of the class.
- ▶ Ask questions, use discussion forums and office hours.

# Key Words In Context (KWIC)

(based on Chapters 8 and 9 in the textbook)

# What is KWIC?

- ▶ KWIC is a way of searching through a text where we do not end our question at: "Where is word X appearing?" but ask for more: "What is the context in which X is appearing?"
- ▶ What does context mean?

# What is KWIC?

- ▶ KWIC is a way of searching through a text where we do not end our question at: "Where is word X appearing?" but ask for more: "What is the context in which X is appearing?"
- ▶ What does context mean?
- ▶ "Context" just refers to the words surrounding X. e.g., two words before and three words after X.
- ▶ The analysis methods we looked at so far saw words as individual entities without bothering about their context.
- ▶ This week, we will study how to extract this context information in R.

# Before we move on ...

What is this function doing?

```
processfile <- function(file_name)
{
  text <- scan(file_name, what = "character", sep = "\n")
  text_as_string <- tolower(paste(text, collapse = " "))
  text_as_string <- gsub("([[:punct:]])", " ", text_as_string)
  text_as_string <- gsub(" +", " ", text_as_string)
  return(unlist(strsplit(text_as_string, " ")))
}
```

What will `processfile("somefilepath")` give me?

# Continuing from last slide:

Try to understand what is happening in this:

```
words <- processfile("DollsHouse-Eng.txt")
helmer <- which(words == "helmer")
for(i in 1:length(helmer))
{
  start <- helmer[i]-2
  end <- helmer[i]+2
  cat(paste(words[start:end], sep=" "))
  cat("\n")
  #What does cat do??
}
```

## Converting the previous slide into a function

```
getKwic <- function(wordsvector, word)
{
  word_index <- which(wordsvector == word)
  for(i in 1:length(word_index))
  {
    start <- word_index[i]-2
    end <- word_index[i]+2
    cat(paste(wordsvector[start:end], sep=" "))
    cat("\n")
  }
}
```

# Interact with the user

- ▶ Why?



# Interact with the user

- ▶ Why?
- ▶ How? - using `readline()` function.

```
file <- readline("Enter the name of the file you want to search in: ")
word <- readline("Enter the word for which you want the context: ")
context0 <- readline("Enter the number of words before and after
                     you want to print: ")
context <- strtoi(readline("Enter the number of words before
                          and after you want to print: "))
```

- ▶ What will happen now?
- ▶ What is the difference between `context0` and `context`?

## Modified getKwic() function

- ▶ The previous getKwic() function we discussed printed out words with a context window of 2 words.
- ▶ It has to be changed to take context from the user input.

```
getKwic <- function(wordsvector, word, context)
{
  word_index <- which(wordsvector == word)
  for(i in 1:length(word_index))
  {
    start <- word_index[i] - context
    end <- word_index[i] + context
    cat(paste(wordsvector[start:end], sep=" "))
    cat("\n")
  }
}
```

# Everything looks good so far

... but

- ▶ What if the the word we are searching for is the very first word?

# Everything looks good so far

... but

- ▶ What if the the word we are searching for is the very first word?
- ▶ a small detour: if I have a vector:  
 $words < -c("Robert", "Rose", "Ryan", "Richard")$
- ▶ What are:

`words[1]`

`words[2]`

`words[0]`

`words[5]`

`words[-1]`

`words[-2]`

`words[1:3]`

`words[-2:3]`

`words[1:7]`

with that knowledge ...


- ▶ Let us come back to the same question: What if the the word we are searching for is the very first word?

with that knowledge ...

- ▶ Let us come back to the same question: What if the the word we are searching for is the very first word?
- ▶ If I am looking for the first word, and am looking for a context of say 3 words, I will get an error at the very beginning, and the R programs stops.
- ▶ If I am looking for the last word, and am looking for a context of say 3 words. What happens???
- ▶ What should we do to avoid these situations??

## final getKwic() function

```
getKwic <- function(wordsvector, word, context)
{
  word_index <- which(wordsvector == word)
  for(i in 1:length(word_index))
  {
    start <- word_index[i] - context
    if(start < 1)
    {
      start <- 1
    }
    end <- word_index[i] + context
    if(end >= length(wordsvector))
    {
      end <- length(wordsvector)
    }
    # cat(start, end) #This prints only positions, not actual words.
    cat(paste(wordsvector[start:end], sep=" "))
    cat("\n")
  }
}
```

Note: Final R file for today (KWIC.R) and will be uploaded to Canvas. 

# How to continue from here?

- ▶ Improve the pre-processing
- ▶ Accept a different context length for before and after (e.g., 2 words before, 3 words after instead of 2 words before and 2 words after)
- ▶ Make this work with more than one file
- ▶ Add the functionality of taking the path to a folder, reading all .txt files from that folder, and do this for each file.
- ▶ Take a doc or pdf, extract plain-text from that and do this
- ▶ Exercises 8.1, 8.2 and 9.1, 9.2 in the textbook (solutions are available in the supplementary material).

.... and so on. (You will do some of this on thursday)



# Ngram analysis

# What are Ngrams?

- ▶ an n-gram is a ordered sequence of words. n- refers to the number of words in the sequence.
- ▶ unigrams - single words, bigrams - two word sequences, trigrams - three word sequences, 4grams - four word sequences and so on.
- ▶ If I have this sentence: "This is an example sentence", what are all possible ngrams in this??

# What are Ngrams?

- ▶ an n-gram is a ordered sequence of words. n- refers to the number of words in the sequence.
- ▶ unigrams - single words, bigrams - two word sequences, trigrams - three word sequences, 4grams - four word sequences and so on.
- ▶ If I have this sentence: "This is an example sentence", what are all possible ngrams in this??
- ▶ Until now, we looked at only words (their frequencies, their position of appearance, their context of appearance)
- ▶ n-gram analysis is about moving beyond words and looking for patterns of word sequences.

# ngram package in R

- ▶ This package provides a lot of functions to automatically create and analyse ngrams from text strings (STRINGS).
- ▶ The package has several advanced functionalities, we don't need at this point.
- ▶ We will only talk about how to extract ngrams of varying sizes from the text, count their frequencies etc.
- ▶ installation: install "ngram" package following the usual procedure.
- ▶ Enthusiastic students can have a look at the documentation for this package to know about all functionalities it has.

## working with ngram package

```
library(ngram)
#usual pre-processing for the file first.
dollshouse_text <- scan("DollsHouse-Eng.txt", what = "character", sep = "\n")
dollshouse_string <- paste(dollshouse_text, collapse = " ")
dollshouse_string <- tolower(dollshouse_string)
dollshouse_string <- gsub("[[:punct:]]", " ", dollshouse_string)
dollshouse_string <- gsub(" +", " ", dollshouse_string)

#Three most useful functions for us:
trigrams <- ngram(dollshouse_string,n=3)
trigrams_vector <- get.ngrams(trigrams)
head(get.phrasetable(trigrams),n = 10)
```

note: Ngram.R file is on Canvas.

## A detour into vectors (again!)

- ▶ R has functions such as `union()` and `intersect()` which takes two vectors and returns:
  - ▶ `union(vector1, vector2)` returns a vector that has all items that occurred in either `vector1` or `vector2`.
  - ▶ `intersection(vector1, vector2)` returns a vector that has all items that occurred in BOTH the vectors.

## Union-Intersection example

```
a = c(1,4,44,5,12)
b = c(1,3,5,44,2)
c <- intersect(a,b)
d <- union(a,b)
c
[1] 1 44 5
d
[1] 1 4 44 5 12 3 2
```

Note: These can be string vectors as well.

Hint: You can use one of these functions to get part of the answer in Question 2 of Assignment 3.

# Assignment 3 Description

- ▶ Topics: Last week, and This Week's content (Chapters 6–9 in Textbook)
- ▶ Questions: 2 questions (4% + 6% of your grade)
- ▶ Deadline: 24th February
- ▶ Description: On Canvas



# Thursday

- ▶ Assignment 2 discussion
- ▶ Practice exercises for using what we learned today (KWIC and N-grams)
- ▶ To do: Read this article "Data Mining reveals the rise of ISIS propaganda on Twitter" (<https://goo.gl/QqTT9k>). We will start the class with a discussion on that on thursday